

```
import pandas as pd
import numpy as np
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```
df = pd.read_csv('/content/drive/MyDrive/concept of Technology of AI/student (1).csv')
df.head()
```

	Math	Reading	Writing
0	48	68	63
1	62	81	72
2	79	80	78
3	76	83	79
4	59	64	62

```
math = df['Math'].to_numpy()
read = df['Reading'].to_numpy()
write = df['Writing'].to_numpy()
```

```
x0 = np.ones(len(math))
X2 = np.array([x0, math, read]).T
W = np.array([0, 0, 0])
Y2 = np.array(write)
```

```
def cost_function(X, Y, W):
    """
    Parameters:
    This function finds the Mean Square Error.
    Input parameters:
        X: Feature Matrix
        Y: Target Matrix
        W: Weight Matrix
    Output Parameters:
        J: accumulated mean square error.
    """
    m = len(Y)

    J = np.sum((X.dot(W) - Y) ** 2)/(2 * m)
    return J
```

```
# Test case
X_test = np.array([[1, 2], [3, 4], [5, 6]])
Y_test = np.array([3, 7, 11])
W_test = np.array([1, 1])
```

```
cost = cost_function(X_test, Y_test, W_test)
```

```
if cost == 0:
    print("Proceed Further")
else:
    print("Something went wrong: Reimplement the cost function")
    print("Cost function output:", cost)
```

Proceed Further

```
initial_cost = cost_function(X2, Y2, W)
print(initial_cost)
```

2470.11

```
def gradient_descent(X, Y, B, alpha, iterations):
    cost_history = [0] * iterations
    m = len(Y)

    for iteration in range(iterations):

        Y_pred = X.dot(B)

        loss = Y_pred - Y

        dw = (X.T.dot(loss)) / (m)

        W_update = W - alpha * dw

        cost = cost_function(X, Y, W_update)
        cost_history[iteration] = cost

    return W_update, cost_history
```

```
alpha = 0.0001
new_weights, cost_history = gradient_descent(X2, Y2, W, alpha, 100000)

print(new_weights)

print(cost_history[-1])
```

[0.0068616 0.4801185 0.5006844]
18.076027940915417

```
def rmse(Y, Y_pred):
    """
    This Function calculates the Root Mean Squares.
    Input Arguments:
        Y: Array of actual(Target) Dependent Variables.
        Y_pred: Array of predeicted Dependent Varaibles.
    Output Arguments:
        rmse: Root Mean Square.
    """
    rmse = np.sqrt(sum((Y-Y_pred)**2)/len(Y))
    return rmse
```

```
def r2(Y, Y_pred):
    """
    This Function calculates the R Squared Error.
    Input Arguments:
        Y: Array of actual(Target) Dependent Varaibles.
        Y_pred: Array of predeicted Dependent Varaibles.
    Output Arguments:
        rsquared: R Squared Error.
    """
    mean_y = np.mean(Y)
    ss_tot = sum((Y - mean_y) ** 2)
    ss_res = sum((Y - Y_pred) ** 2)
    r2 = 1 - (ss_res / ss_tot)
    return r2
```

```
Y_pred = X2.dot(new_weights)

print(rmse(Y2, Y_pred))
print(r2(Y2, Y_pred))
```

6.01265797146577
0.8442155132417348

```
import pandas as pd
import numpy as np
```

- To - Do - 1:

Read and Observe the Dataset. Print top(5) and bottom(5) of the dataset {Hint: pd.head and pd.tail}. Print the Information of Datasets. {Hint: pd.info}. Gather the Descriptive info about the Dataset. {Hint: pd.describe} Split your data into Feature (X) and Label (Y).

```
df = pd.read_csv('/content/drive/MyDrive/concept of Technology of AI/Copy of Houseprice (1).csv');
print(df.head())
df.tail()
```

	HouseAge	HouseFloor	HouseArea	HousePrice
0	52	2	112.945574	543917.179841
1	93	1	174.312126	817740.124828
2	15	4	125.219577	387992.503019
3	72	4	121.210124	240840.742388
4	61	4	59.221737	277273.386525
	HouseAge	HouseFloor	HouseArea	HousePrice
95	85	1	120.193091	540347.535683
96	80	1	56.078992	222539.252789
97	82	2	211.368074	715211.774907
98	53	4	94.277670	210926.069798
99	24	4	285.114646	988329.950912

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 4 columns):
 #   Column      Non-Null Count  Dtype  
 ---  --          --          --      
 0   HouseAge    100 non-null   int64  
 1   HouseFloor  100 non-null   int64  
 2   HouseArea   100 non-null   float64 
 3   HousePrice  100 non-null   float64 
dtypes: float64(2), int64(2)
memory usage: 3.3 KB
```

```
df.isnull().sum() #NO empty data
```

	0
HouseAge	0
HouseFloor	0
HouseArea	0
HousePrice	0

```
dtype: int64
```

```
df.describe()
```

	HouseAge	HouseFloor	HouseArea	HousePrice
count	100.000000	100.000000	100.000000	1.000000e+02
mean	51.540000	2.910000	170.226332	5.790964e+05
std	29.425963	1.457097	71.298821	2.615097e+05
min	2.000000	1.000000	51.265396	1.493562e+05
25%	23.500000	1.000000	109.770201	3.905859e+05
50%	54.000000	3.000000	176.846510	5.391941e+05
75%	76.500000	4.000000	222.794276	7.479899e+05
max	100.000000	5.000000	296.412614	1.228712e+06

```
Age = df["HouseAge"].to_numpy().ravel()
Floor = df["HouseFloor"].to_numpy().ravel()
Area = df["HouseArea"].to_numpy().ravel()
Price = df["HousePrice"].to_numpy().ravel()
```

```
X = df[['HouseAge', 'HouseFloor', 'HouseArea']] # features
Y = df['HousePrice'] #Label
W = np.array([-1.5, 4, 6])
```

```
def scale_features(X):
    """Scale features (excluding intercept column)"""
    X_scaled = X.copy()
    for col in range(1, X.shape[1]): # Skip intercept column (column 0)
        mean = np.mean(X[:, col])
        std = np.std(X[:, col])
        if std > 0: # Avoid division by zero
            X_scaled[:, col] = (X[:, col] - mean) / std
        else:
            X_scaled[:, col] = 0
    return X_scaled
```

```
def scale_target(Y):
    """Scale target variable"""
    mean = np.mean(Y)
    std = np.std(Y)
    Y_scaled = (Y - mean) / std
    return Y_scaled, mean, std
```

TODO 2:

```
W = np.array([1, 1, 1]).reshape(-1,1)
print(W.shape)
print(X.shape)
result= np.dot(X,W)
print(result[:5])

(3, 1)
(100, 3)
[[166.94557396]
 [268.31212647]
 [144.21957745]
 [197.21012359]
 [124.22173684]]
```

TODO 3:

```
def split(X, Y, test_split=0.2, seed=42):
    np.random.seed(seed)
    indices = np.arange(X.shape[0])
    np.random.shuffle(indices)

    split_size = int(len(X) * test_split)

    test_indices = indices[:split_size]
    train_indices = indices[split_size:]

    X_train, X_test = X[train_indices], X[test_indices]
    Y_train, Y_test = Y[train_indices], Y[test_indices]

    return X_train, X_test, Y_train, Y_test
```

TODO 4:

```
x0 = np.ones(len(Age))
X4 = np.array([x0, Age, Floor, Area]).T
W4 = np.array([0.0, 0.0, 0.0, 0.0])
Y4 = Price.ravel()
```

```
X4_scaled = scale_features(X4)
Y4_scaled, Y_mean, Y_std = scale_target(Y4)
```

```
def cost_function(X, Y, W):
    """ Parameters:
        This function finds the Mean Square Error.
        Input parameters:
            X: Feature Matrix
            Y: Target Matrix
            W: Weight Matrix
        Output Parameters:
            J: accumulated mean square error.
    """
    m = len(Y)

    J = np.sum((X.dot(W) - Y) ** 2)/(2 * m)
    return J
```

```
# Test case
X_test = np.array([[1, 2],
                  [3, 4],
                  [5, 6]])  #3X2
Y_test = np.array([3, 7, 11])  #1X3
W_test = np.array([1, 1])  # 1X2

cost = cost_function(X_test, Y_test, W_test)

if cost == 0:
    print("Proceed Further")
else:
    print("Something went wrong: Reimplement the cost function")
    print("Cost function output:", cost)
```

Proceed Further

```
initial_cost = cost_function(X4, Y4, W4)
print(initial_cost)
```

201528080199.132

```
print(X4.shape)
print(Y4.shape)
print(W4.shape)

(100, 4)
(100,)
(4,)
```

```
def gradient_descent(X, Y, B, alpha, iterations):
    cost_history = [0] * iterations
    m = len(Y)

    for iteration in range(iterations):
        # Forward pass
        Y_pred = X.dot(B)
        loss = Y_pred - Y

        # Gradient calculation
        gradient = (X.T.dot(loss)) / m

        # Update weights
        B = B - alpha * gradient  # FIX: Update B, not create W_update

        # Calculate and store cost
        cost_history[iteration] = cost_function(X, Y, B)

    return B, cost_history
```

```
def rmse(Y, Y_pred):
    """
    This Function calculates the Root Mean Squares.
    Input Arguments:
```

```

Y: Array of actual(Target) Dependent Variables.
Y_pred: Array of predicted Dependent Variables.
Output Arguments:
    rmse: Root Mean Square.
"""
rmse = np.sqrt(sum((Y-Y_pred)**2)/len(Y))
return rmse

```

```

def r2(Y, Y_pred):
    """
    This Function calculates the R Squared Error.
    Input Arguments:
        Y: Array of actual(Target) Dependent Variables.
        Y_pred: Array of predicted Dependent Variables.
    Output Arguments:
        rsquared: R Squared Error.
    """
    mean_y = np.mean(Y)
    ss_tot = sum((Y - mean_y) ** 2)
    ss_res = sum((Y - Y_pred) ** 2)
    r2 = 1 - (ss_res / ss_tot)
    return r2

```

```

def simple_normalize(arr):
    """Normalize array to [0, 1] range"""
    return (arr - np.min(arr)) / (np.max(arr) - np.min(arr))

X4_norm = X4.copy()
Y4_norm = simple_normalize(Y4)

```

```

for i in range(1, X4_norm.shape[1]):
    X4_norm[:, i] = simple_normalize(X4[:, i])

print("== DATA NORMALIZED TO [0, 1] RANGE ==")
print(f"X ranges: [{np.min(X4_norm[:, 1]):.3f}, {np.max(X4_norm[:, 1]):.3f}]")
print(f"Y range: [{np.min(Y4_norm):.3f}, {np.max(Y4_norm):.3f}]")

== DATA NORMALIZED TO [0, 1] RANGE ==
X ranges: [0.000, 1.000]
Y range: [0.000, 1.000]

```

```

initial_cost = cost_function(X4, Y4, W4)
print(f"Initial cost: {initial_cost:.2f}")

Initial cost: 201528080199.13

```

```

alpha = 0.000025
new_weights, cost_history = gradient_descent(X4, Y4, W4, alpha, 10000)

print(f"\nOptimized weights:")
print(f"Intercept (b0): {new_weights[0]:.4f}")
print(f"HouseAge coefficient (b1): {new_weights[1]:.4f}")
print(f"HouseFloor coefficient (b2): {new_weights[2]:.4f}")
print(f"HouseArea coefficient (b3): {new_weights[3]:.4f}")
print(f"Final cost: {cost_history[-1]:.2f}")

```

```

Optimized weights:
Intercept (b0): 1982.9551
HouseAge coefficient (b1): 472.6966
HouseFloor coefficient (b2): -4256.0767
HouseArea coefficient (b3): 3270.8719
Final cost: 12016392087.97

```

```

Y_pred = X4.dot(new_weights)

print(f"\nModel Performance:")
print(f"RMSE: {rmse(Y4, Y_pred):.2f}")
print(f"R-squared: {r2(Y4, Y_pred):.4f}")

```

```

Model Performance:
RMSE: 155025.11

```

R-squared: 0.6450

```
print("\nFirst 5 predictions vs actual:")
for i in range(5):
    print(f"Actual: {Y4[i]:.2f}, Predicted: {Y_pred[i]:.2f}, Difference: {abs(Y4[i] - Y_pred[i]):.2f}")
```

```
First 5 predictions vs actual:
Actual: 543917.18, Predicted: 387481.53, Difference: 156435.65
Actual: 817740.12, Predicted: 611840.30, Difference: 205899.82
Actual: 387992.50, Predicted: 401626.30, Difference: 13633.79
Actual: 240840.74, Predicted: 415455.59, Difference: 174614.85
Actual: 277273.39, Predicted: 207499.86, Difference: 69773.53
```

```
new_weights_scaled, cost_history = gradient_descent(
    X4_scaled, Y4_scaled, W4, alpha, 10000
)
```

```
print(f"\nModel Performance:")
print(f"RMSE: {rmse(Y4, Y_pred):.2f}")
print(f"R-squared: {r2(Y4, Y_pred):.4f}")
```

```
Model Performance:
RMSE: 155025.11
R-squared: 0.6450
```

```
print("\nFirst 5 predictions vs actual:")
for i in range(5):
    print(f"Actual: {Y4[i]:.2f}, Predicted: {Y_pred[i]:.2f}, Difference: {abs(Y4[i] - Y_pred[i]):.2f}")

# Also show cost reduction
print("\nCost reduction:")
print(f"Initial cost: {cost_history[0]:.4f}")
print(f"Final cost: {cost_history[-1]:.4f}")
```

```
First 5 predictions vs actual:
Actual: 543917.18, Predicted: 387481.53, Difference: 156435.65
Actual: 817740.12, Predicted: 611840.30, Difference: 205899.82
Actual: 387992.50, Predicted: 401626.30, Difference: 13633.79
Actual: 240840.74, Predicted: 415455.59, Difference: 174614.85
Actual: 277273.39, Predicted: 207499.86, Difference: 69773.53
```

```
Cost reduction:
Initial cost: 0.5000
Final cost: 0.3663
```