

```
import pandas as pd
import numpy as np
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
```

2 Exercise 1 - Implementation of Naive Bayes Algorithm

Question: You are tasked with building a sentiment analysis model to classify movie reviews as positive or negative using the IMDB Movie Review Dataset.

Part 1:

1. Load the dataset and preprocess the reviews.
 - a. Convert all text to lowercase.
 - b. Remove non-alphabetic characters (punctuation).
 - c. Tokenize the reviews and remove common stopwords.
 - d. Apply stemming to reduce words to their root form.

```
nltk.download('stopwords')
df = pd.read_csv("/content/drive/MyDrive/ConceptAndTechnologyOfAI/IMDB.csv")
ps = PorterStemmer()
stop_words = set(stopwords.words('english'))

def preprocess_text(text):
    text = text.lower()
    text = re.sub('[^a-z]', ' ', text)
    tokens = text.split()
    tokens = [ps.stem(word) for word in tokens if word not in stop_words]
    return " ".join(tokens)

df['clean_review'] = df['review'].apply(preprocess_text)

df['label'] = df['sentiment'].map({'positive':1, 'negative':0})
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
-----
FileNotFoundError                                Traceback (most recent call last)
/tmp/ipython-input-3687145463.py in <cell line: 0>()
      1 nltk.download('stopwords')
----> 2 df = pd.read_csv("/content/drive/MyDrive/Concept and technology of AI/IMDB.csv")
      3 ps = PorterStemmer()
      4 stop_words = set(stopwords.words('english'))
      5

-----
4 frames -----
/usr/local/lib/python3.12/dist-packages/pandas/io/common.py in get_handle(path_or_buf, mode, encoding, compression, memory_map,
is_text, errors, storage_options)
    871         if ioargs.encoding and "b" not in ioargs.mode:
    872             # Encoding
--> 873             handle = open(
    874                 handle,
    875                 ioargs.mode,
FileNotFoundError: [Errno 2] No such file or directory: '/content/drive/MyDrive/Concept and technology of AI/IMDB.csv'
```

Next steps: [Explain error](#)

```
from google.colab import drive
drive.mount('/content/drive')
```

2. Split the dataset into training and testing sets (80% training, 20% testing).

```
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import MultinomialNB
```

```
X_train, X_test, y_train, y_test = train_test_split(df['clean_review'], df['label'], test_size=0.2, random_state=42)
```

3. Use a Naive Bayes classifier to classify the reviews into positive and negative categories.

a. Implement a Bag-of-Words model using CountVectorizer.

b. Train the Naive Bayes classifier using the training set.

```
vectorizer = CountVectorizer(max_features=5000)
X_train_vec = vectorizer.fit_transform(X_train).toarray()
X_test_vec = vectorizer.transform(X_test).toarray()

nb = MultinomialNB()
nb.fit(X_train_vec, y_train)

y_pred = nb.predict(X_test_vec)
```

Part 2:

1. Evaluate the performance of the model using the following metrics:

a. Accuracy

b. Precision, Recall, and F1-score

c. Confusion Matrix

d. ROC-AUC Score

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, roc_auc_score, classification_report
import seaborn as sns
import matplotlib.pyplot as plt

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

3 Feature Selection Exercise.

3.1 Feature Selection using Wrapper Methods.

Dataset: Breast Cancer Prognostic Dataset

Objective You are tasked with identifying the most important features in predicting breast cancer prognosis using the Breast Cancer Prognostic Dataset from the UCI Machine Learning Repository. Apply a wrapper method (Recursive Feature Elimination) to select the best features and evaluate your model's performance.

Part 1: Data Loading and Preprocessing

1. Load the Breast Cancer Prognostic Dataset.
2. Dataset is available in Drive.
3. Perform basic exploratory data analysis (EDA) to understand the dataset:
 - Summarize key statistics for each feature.
 - Check for missing values and handle them appropriately.
4. Split the dataset into training (80%) and testing (20%) sets.

```
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
from sklearn.model_selection import train_test_split
```

```

data = pd.read_csv("/content/drive/MyDrive/Concept and technology of AI/BreastCancer.csv")

print(data.info())
print(data.describe())
print(data.isnull().sum())

numeric_cols = data.select_dtypes(include=['float64', 'int64']).columns
categorical_cols = data.select_dtypes(include=['object']).columns

data[numeric_cols] = data[numeric_cols].fillna(data[numeric_cols].median())
for col in categorical_cols:
    data[col] = data[col].fillna(data[col].mode()[0])

for col in categorical_cols:
    data[col] = data[col].astype('category').cat.codes

print(data.columns)
X = data.drop('diagnosis', axis=1)
y = data['diagnosis'].astype('category').cat.codes

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

Part 2: Apply a Wrapper Method

1. Use Recursive Feature Elimination (RFE) with a Logistic Regression model to perform feature selection:

- Select the top 5 features that contribute the most to predicting the target variable.
- Visualize the ranking of features.

2. Train the Logistic Regression model using only the selected features.

```

from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
import pandas as pd
import matplotlib.pyplot as plt

imputer = SimpleImputer(strategy='median')
X_train_imp = imputer.fit_transform(X_train)
X_test_imp = imputer.transform(X_test)

feature_names = X.columns

if X_train_imp.shape[1] != len(feature_names):
    print("Column mismatch detected!")
    print("X_train shape:", X_train_imp.shape)
    print("Feature names count:", len(feature_names))
else:
    X_train = pd.DataFrame(X_train_imp, columns=feature_names)
    X_test = pd.DataFrame(X_test_imp, columns=feature_names)

    logreg = LogisticRegression(max_iter=1000)
    rfe = RFE(logreg, n_features_to_select=5)
    rfe.fit(X_train, y_train)

    selected_features = X_train.columns[rfe.support_]
    print("Top 5 Selected Features:", selected_features)

    ranking = pd.DataFrame({'Feature': X_train.columns, 'Rank': rfe.ranking_})
    ranking.sort_values('Rank', inplace=True)

    plt.barh(ranking['Feature'], ranking['Rank'])
    plt.xlabel("RFE Rank (1=Best)")
    plt.title("Feature Ranking")
    plt.show()

```

Part 3: Model Evaluation

1. Evaluate the model's performance using the testing set:

- Metrics to calculate: Accuracy, Precision, Recall, F1-Score, and ROC-AUC.

2. Compare the performance of the model trained on all features versus the model trained on the selected features.

```
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.feature_selection import RFE
import pandas as pd
import matplotlib.pyplot as plt

imputer = SimpleImputer(strategy='median')
X_train_imp = imputer.fit_transform(X_train)
X_test_imp = imputer.transform(X_test)

feature_names = X.columns[:X_train_imp.shape[1]]

X_train = pd.DataFrame(X_train_imp, columns=feature_names)
X_test = pd.DataFrame(X_test_imp, columns=feature_names)

logreg = LogisticRegression(max_iter=1000)
rfe = RFE(logreg, n_features_to_select=5)
rfe.fit(X_train, y_train)

selected_features = X_train.columns[rfe.support_]
print("Top 5 Selected Features:", selected_features)

ranking = pd.DataFrame({'Feature': X_train.columns, 'Rank': rfe.ranking_})
ranking.sort_values('Rank', inplace=True)

plt.barh(ranking['Feature'], ranking['Rank'])
plt.xlabel("RFE Rank (1=Best)")
plt.title("Feature Ranking")
plt.show()
```

Part 4: Experiment

1. Experiment with different numbers of selected features (e.g., top 3, top 7).
2. Discuss how feature selection affects model performance.

```
for n in [3, 7]:
    rfe = RFE(logreg, n_features_to_select=n)
    rfe.fit(X_train, y_train)
    selected = X_train.columns[rfe.support_]
    print(f"\nTop {n} Features:", selected.tolist())

    logreg.fit(X_train[selected], y_train)
    y_pred_n = logreg.predict(X_test[selected])
    print(f"Performance with Top {n} Features:")
    print("Accuracy:", accuracy_score(y_test, y_pred_n))
    print("F1 Score:", f1_score(y_test, y_pred_n))
```